

Handout 4: Exchangeability and the Bayesian model

Lecturer & author: Georgios P. Karagiannis

georgios.karagiannis@durham.ac.uk

Aim

Get familiar with the concept of exchangeability, and its relation to Subjective probability and Bayesian paradigm.

Reading list:

- Bernardo, J. M., & Smith, A. F. (2009, Section 4.3). Bayesian theory (Vol. 405). John Wiley & Sons.
- Berger, J. O. (2013, Section 3.5.7). Statistical decision theory and Bayesian analysis. Springer Science & Business Media.

1 Exchangeability

As mentioned in Handout 3, one way to specify the Bayesian model is by subjectively specifying the probability distributions $F(y|\theta)$ and $\Pi(\theta)$, or the joint distribution $P(y, \theta)$, that enables You to derive the rest distributions.

Alternatively, You can specify a probability distribution on the data generating process $G(y_{1:n})$ describing the actual sequence of the data $y_{1:n} = (y_1, \dots, y_n)$. One approach is to subjectively set certain invariance assumptions on the observables y involving probabilistic believes of invariant with respect to some aspect of the observable quantities.

A reasonable invariance assumption about $y_{1:n} = (y_1, \dots, y_n)$ (to specify the data generating model $G(y_{1:n})$) is the Exchangeability: The ‘labels’ identifying the individual observable quantities are ‘uninformative’, in the sense the information that the y_i ’s provide is independent of the order in which they are collected. Exchangeability, although a simple assumption, it accurately describes a large class of experimental setups.

Definition 1. A sequence of random quantities $y_{1:n} = (y_1, \dots, y_n)$ is finitely exchangeable under a probability distribution G if all permutations of $\{y_1, \dots, y_n\}$ have the same joint distribution G . Namely; if

$$G(y_1, \dots, y_n) = G(y_{p(1)}, \dots, y_{p(n)})$$

for all permutations p defined on the set $\{1, \dots, n\}$.

Definition 2. An infinite sequence of random quantities y_1, y_2, \dots is infinitely exchangeable under a probability distribution G if every finite sub-sequence is finitely exchangeable under G .

Example 3. If a sequence of random quantities $y_{1:n} = (y_1, \dots, y_n)$ is mutually independent, then it is exchangeable.

Solution. It is straightforward, since $G(y_1, \dots, y_n) = \prod_{i=1}^n G(y_i)$ which is invariant to permutations of the indexes.

2 The representation theorem

Note 4. The assumption of exchangeability leads to the following development, which theoretically justifies (to some extend) the existence of the Prior distribution, and the Bayesian paradigm.

Theorem 5. (General representation theorem) If y_1, y_2, \dots is an infinitely exchangeable sequence of random quantities with probability distribution F , there exists a probability measure Π over \mathcal{F} , the space of all distribution functions on

$\mathcal{Y}^n \subseteq \mathbb{R}^n$ for $n \geq 1$, such that the joint distribution function of $y_{1:n} = (y_1, \dots, y_n)$ has CDF

$$G(y_1, \dots, y_n) = \int_{\mathcal{F}} \prod_{i=1}^n F(y_i) d\Pi(F) \quad (1)$$

where F is an unknown/unobservable distribution function, $\Pi(F) = \lim_{n \rightarrow \infty} P(\hat{F}_n)$ is a probability distribution on the space of functions \mathcal{F} , which is defined as a limit distribution on the empirical distribution function $\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n 1(y_i \leq x)$ defined by y_1, \dots, y_n (as $n \rightarrow \infty$), and $F(x) = \lim_{n \rightarrow \infty} \hat{F}_n(x)$.

Remark 6. The Representation Theorem shows that if y_1, y_2, \dots is infinitely exchangeable, then the elements of $y_{1:n} = (y_1, \dots, y_n)$ are i.i.d. conditional on the empirical distribution of $y_{1:n}$.

Remark 7. (Interpretation) The general representation Theorem 5 says

- $y_{1:n}$ are considered to be an i.i.d. sample generated from an unknown (i.e., random) distribution function F , (conditional on F); i.e. $y_i|F \sim F(\cdot)$.
- F is an the unknown CDF, which follows itself a probability distribution Π representing (prior) believes about F
- and F has the operational role of what You believe the empirical distribution function would look like for a large sample.

Note 8. For convenience, hereafter, we will consider the concept of exchangeability by using the parametric form 2 and 3 in Fact 9.

Fact 9. Given some additional problem specific invariance assumptions (see Bernardo & Smith (2009, Section 4.3)) via subjunctive judgments regarding the generating process of y_1, y_2, \dots , the unknown sampling distribution $G(y_i)$ in (1) can be written as a parametric model $F(y_i|\theta)$ labeled by an unknown parameter $\theta \in \Theta$, which is the limit of some function of $y_{1:n}$ (as $n \rightarrow \infty$), and there exists a probability distribution Π for θ such that

$$G(y_1, \dots, y_n) = \int_{\Theta} \prod_{i=1}^n F(y_i|\theta) d\Pi(\theta) \quad (2)$$

In PDF/PMF, (2) is written as

$$g(y_1, \dots, y_n) = \int_{\Theta} \prod_{i=1}^n f(y_i|\theta) d\Pi(\theta) \quad (3)$$

Remark 10. Regarding the Bayesian paradigm, Fact 9 provides a rational for the consideration of the uncertain parameter θ as a random variable and the subjective prior $\Pi(\theta)$. If y_1, y_2, \dots is an exchangeable sequence of real-valued random quantities, then any finite subset of them is an i.i.d. random sample from parametric model $F(\cdot|\theta)$ labeled by some uncertain parameter $\theta \in \Theta$, and there exists a (prior) probability distribution $\Pi(\theta)$ for θ which has to describe the initially available information about the parameter which labels the model; Hence a rational for the parametric Bayesian model:

$$\begin{cases} y_i|\theta & \stackrel{\text{iid}}{\sim} F(\cdot|\theta) \quad \forall i = 1, \dots, n \\ \theta & \sim \Pi(\cdot) \end{cases}$$

Note 11. The following Example 12 of the ‘representation Theorem with 0 – 1 quantities’ presents a special case.

Example 12. (Representation Theorem with 0 – 1 quantities). If y_1, y_2, \dots is an infinitely exchangeable sequence of 0 – 1 random quantities with probability measure G , there exists a distribution function Π such that the joint mass function $g(y_1, \dots, y_n)$ for y_1, \dots, y_n has the form

$$g(y_1, \dots, y_n) = \int_0^1 \prod_{i=1}^n \underbrace{\theta^{y_i} (1 - \theta)^{1-y_i}}_{f_{\text{Br}(\theta)}(y_i|\theta)} d\Pi(\theta)$$

where

$$\Pi(t) = \lim_{n \rightarrow \infty} P\left(\frac{1}{n} \sum_{i=1}^n y_i \leq t\right) \quad \text{and} \quad \theta \stackrel{\text{as}}{=} \lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n y_i$$

aka θ is the limiting relative frequency of 1s, by SLLN.

Solution. It is given in the Exercise sheet to read; see Exercise 34.

Example 13. (cont. Example 12) The representation of exchangeable sequence of 0 – 1 random quantities y_1, \dots, y_n can be interpreted as follows:

- y_i are considered to be conditionally independent and identically distributed Bernoulli random quantities given the random quantity θ ; i.e. $y_i|\theta \stackrel{iid}{\sim} \text{Br}(\theta)$.
- θ is itself assigned a probability distribution Π which can be interpreted as its prior distribution,
- by the SLLN, θ is defined as $\theta = \lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n y_i$, and hence Π can be interpreted as beliefs about the limiting relative frequency of 1's.

Note 14. The following Example 15 can justify the notion of prior, posterior and predictive distribution in the context of the exchangeability.

Example 15. Let y_1, y_2, \dots be an infinitely exchangeable sequence of random quantities under distribution G admitting a PDF/PMF g . Then from the representation theorem in (2), the conditional distribution $G(y_{n+1:n+m}|y_{1:n})$ has PDF/PMF

$$g(y_{n+1:n+m}|y_{1:n}) = \int_{\Theta} \prod_{i=n+1}^{n+m} f(y_i|\theta) d\Pi(\theta|y_{1:n}) \quad \text{where} \quad d\Pi(\theta|y_{1:n}) = \frac{\prod_{i=1}^n f(y_i|\theta) d\Pi(\theta)}{\int_{\Theta} \prod_{i=1}^n f(y_i|\theta) d\Pi(\theta)}$$

Solution. It can be shown that

$$\begin{aligned} g(y_{n+1:n+m}|y_{1:n}) &= \frac{g(y_{1:n}, y_{n+1:n+m})}{g(y_{1:n})} = \frac{\int_{\Theta} \prod_{i=1}^n f(y_i|\theta) \prod_{i=n+1}^{n+m} f(y_i|\theta) d\Pi(\theta)}{\int_{\Theta} \prod_{i=1}^n f(y_i|\theta) d\Pi(\theta)} \\ &= \int_{\Theta} \prod_{i=n+1}^{n+m} f(y_i|\theta) \underbrace{\frac{\prod_{i=1}^n f(y_i|\theta) d\Pi(\theta)}{\int_{\Theta} \prod_{i=1}^n f(y_i|\theta) d\Pi(\theta)}}_{=d\Pi(\theta|y_{1:n})} \end{aligned}$$

Remark 16. Subjectively specifying $G(y_{1:n})$ and then deriving $F(y_{1:n}|\theta)$ and $\Pi(\theta)$ is philosophically interesting. It can suggest useful sampling-model & prior ($F(y_{1:n}|\theta)$, $\Pi(\theta)$) decompositions, that allow the design of new meaningful models. For example, see Bernardo, J. M., & Smith, A. F. (2009, Section 4.4). On the other hand, it is often easier to subjectively specify the Bayesian model by $F(y_{1:n}|\theta)$ and $\Pi(\theta)$.

Example 17. Consider the parametric form (2) of the general representation theorem. Let y_1, y_2, \dots be an infinitely exchangeable sequence of real valued random quantities with $y_i \in \mathbb{R}$ for any i .

1. Show that $\text{Corr}(y_i, y_j) \geq 0$, for $i \neq j$.
2. Find the condition under which (i.) I have $\text{Corr}(y_i, y_j) > 0$, for $i \neq j$ (ii.) I have $\text{Corr}(y_i, y_j) = 0$, for $i \neq j$

Solution. Since sequence y_1, y_2, \dots is infinitely exchangeable, I use the general representation theorem (the parametric form for simplicity). Hence, for a given θ , it is $x_i|\theta \stackrel{iid}{\sim} dF(\cdot|\theta)$ for all i .

1. So

$$\begin{aligned} \text{Cov}_G(y_i, y_j) &= E_G(x_i^\top x_j) - E_G(x_i)^\top E_G(x_j) = E_G(E_F(x_i^\top x_j|\theta)) - E_G(E_F(x_i|\theta))^\top E_G(E_F(x_j|\theta)) \\ &= E_\Pi(E_F(x_i^\top|\theta)E_F(x_j|\theta)) - E_\Pi(E_F(x_i|\theta))^\top E_\Pi(E_F(x_j|\theta)) = \text{Var}_\Pi(\mu(\theta)) \geq 0 \end{aligned}$$

where $\mu(\theta) = E_F(x_i|\theta)$ for all i . Also

$$\text{Var}_G(y_i) = \text{Var}_\Pi(E_F(x_j|\theta)) + E_\Pi(\text{Var}_F(x_j|\theta)) = \text{Var}_\Pi(\mu(\theta)) + E_\Pi(\sigma^2(\theta))$$

where $\sigma^2(\theta) = \text{Var}_F(x_i|\theta)$ for all i . Lets consider the 1-D case, that is requested; $y_i \in \mathbb{R}$ for any i . It is

$$\text{Corr}(y_i, y_j) = \frac{\text{Var}_\Pi(\mu(\theta))}{\text{Var}_\Pi(\mu(\theta)) + E_\Pi(\sigma^2(\theta))} \geq 0$$

2. For $\text{Var}_\Pi(E_F(x_i|\theta)) > 0$, I have $\text{Corr}(y_i, y_j) > 0$. For $\text{Var}_\Pi(E_F(x_i|\theta)) = 0$, I have $\text{Corr}(y_i, y_j) = 0$;

NB: correlation does not necessarily imply independence.

Remark 18. Example 17 shows that elements of infinite exchangeable sequence cannot be negatively correlated.

3 Practice

Question 19. For practice try the Exercises 36, 37, 38, and 39 from the Exercise Sheet.